

Annotation Guidelines

Contents

1	Objective	2
2	What you're doing	2
2.1	Annotation units	2
2.2	Annotation principles	2
2.3	What not to annotate (exclusion criteria)	3
3	What you'll see for each message	3
4	Language labels (one per token)	3
4.1	Label precedence	3
4.2	Language decision tree (use for each token)	4
4.3	Rules where mistakes happen	4
5	Harm labels (per message)	5
5.1	Q1: is there harmful or offensive content?	5
5.2	Q1a: harmful content in non-English terms	6
5.3	Q2: is it aggressive, and is it real or playful?	6
5.4	Target	7
5.5	Harm decision tree	7
6	Worked examples (the hard cases)	7
7	The order to annotate	10
8	Workflow, disagreement & quality	11
8.1	Confidence: what "uncertain" means	11
8.2	Pilot / calibration round	11
8.3	Independent annotation & adjudication	11
8.4	Inter-annotator agreement	12
8.5	Quality expectations	12

1 Objective

Produce consistent token level language annotations and message level harm annotations for conversational Singaporean code-mixed text. We require two careful annotators following these rules land on the same labels, and that any disagreement stays in the data.

2 What you're doing

The annotator is to read short chat messages from real Singaporean conversations: casual texting between friends that mixes English with Malay, Chinese, Tamil, and Singlish slang. Two kinds of labels are to be added:

1. **Language labels:** for each word, which language it belongs to.
2. **Harm labels:** for each message, whether it contains anything harmful or hostile, and of what kind.

The core rule

Language labels identify what each token is (which language it belongs to) and harm labels identify what the message is doing in context. The same words can mean completely different things depending on how they're used and who is talking. Friends insult each other affectionately (e.g. "you idiot" with a laughing emoji between close friends is not an attack), and a swear word can be venting ("ugh, damn") rather than an attack. So for the harm labels, judge what the message is doing, not just which words appear. Two messages with identical words can get different harm labels.

Annotators are not expected to be fluent in every language in the dataset. You can still annotate consistently when you hit unfamiliar words: use the rules below and the `uncertain` option. If the evidence isn't there, mark `uncertain`.

2.1 Annotation units

- Language identification is at the token level. Each token gets exactly one language label, and the labels are mutually exclusive. There is no "this is both HK and SG" option. If a token feels like it could be two things, resolve it with the rules in Section 4 and pick one.
- Harm annotation is at the message level. Each message receives one harm-type label, one aggression label, one target label, and one confidence label, except excluded turns (see the exclusion criteria below).
- The preceding turns plus the relationship is used to interpret message-level harm annotations. Context is not used when it comes to language identification ONLY for harm annotations.

2.2 Annotation principles

Keep these three rules in mind:

1. Label what is present, not what is implied.
2. Use context for aggression and function, not for language labels.
3. When the evidence isn't there, mark uncertain rather than inferring.

2.3 What not to annotate (exclusion criteria)

Some rows are not annotatable text. System messages, deleted message placeholders, media placeholders (GIF / sticker / image / attachment, and malformed or empty turns do not receive harm labels, and their non-linguistic placeholder tokens are tagged `OTH` (Section 4). A lone reaction emoji (a thumbs-up on its own) is `OTH` and gets no harm label. If a whole turn is a placeholder, skip the harm annotation for that turn and note it.

This corpus was self-anonymised by donors, so you will see bracketed stand-ins. Tag `[NAME]` and `[LOCATION]` as `NE` (they stand in for a named entity). Tag `[DATE]`, `[TIME]`, and similar as `OTH` (they stand in for numeric or temporal content). A placeholder by itself is never harmful content.

3 What you'll see for each message

1. The message, already split into **tokens** (a token = one word, emoji, or punctuation mark). You tag each token and do **not** split anything yourself.
2. The conversation before it, usually the previous 3 turns. This is the *context*. Always read it before the harm labels. If those few turns are just fragments (“haha”, “ok”) and don’t make the situation clear, look a little further back. Read enough to understand what’s going on (usually three turns is enough).
3. The relationship between the two people (“close friend”, “colleague”, “family”). Rough teasing between close friends usually bares no ill intent based on the relationship between the people. However, the same words between coworkers may not be.

For some conversations, we don’t know who sent which message. In that case, lean toward `uncertain` if not knowing the speaker leaves the call unclear.

4 Language labels (one per token)

Assign exactly one label to each token, in order. Do not merge or split them. If a token looks wrongly split (e.g. a number stuck to a word), tag it as best you can and flag it.

4.1 Label precedence

Two of the labels are **overrides**. Apply them before anything else, whatever the word looks like:

- `NE`: a name, place, or brand is `NE` even when it is also a common word (e.g. “Apple” the company, not the fruit).
- `OTH`: emoji, URLs, numbers, punctuation, and laughter are `OTH` even though laughter (“HAHAHA”) is alphabetic.

After the two overrides, choosing among `EN`, `SG`, and the other languages is a judgement about how the token is used in this message. There is no priority order. In particular, `SG` does not automatically beat English: a token is `SG` only when it is actually functioning as a Singlish colloquial item in this message (see the disambiguation rules below). Use `UNK` only when you really cannot place a token in any language.

Label	Use it for	Examples
EN	English, including informal spellings, texting shortcuts, and internet slang	u, dat, tgt, alr, btr, idk, lol, lmao, wtf, wanna
ZH	Chinese (Mandarin): Chinese characters or romanised Chinese (pinyin)	睡觉, 我, 不知道, “ni hao”, “xie xie”
MS	Malay words (including romanised)	makan, tak, bisa
TA	Tamil words (<i>rare in this data; including romanised</i>)	(Tamil-script or romanised Tamil)
HK	Hokkien words (<i>rare in this data</i>)	(certain dialect words)
SG	Singlish colloquial items: particles <i>and</i> borrowed words used as Singlish	lah, leh, lor, sia, meh, hor, liao, eh, wah; paiseh, kiasu, sian
NE	Names of people, places, brands, apps	Annette, Suntec, ChatGPT, Salomon, [NAME], [LOCATION]
OTH	Emoji, web links, numbers, punctuation, and laughter (even alphabetic)	emoji, https://..., 123, !!!, HAHABA, hahaha, [TIME]
UNK	You genuinely cannot tell what it is	(rare; only when truly stuck)

A note on why **SG** groups particles and borrowings: annotators can’t be expected to know that *paiseh* is Hokkien derived while *lah* is Chinese derived. What matters is whether the token is a colloquial Singlish item as used.

4.2 Language decision tree (use for each token)

Work down the list and stop at the first match

1. Is it a name / place / brand / app ([NAME], [LOCATION])? → **NE**
2. Is it emoji / link / number / punctuation / laughter / [TIME]? → **OTH**
3. Is it functioning as a Singlish colloquial item *in this message* (lah, sia, paiseh, kiasu)? → **SG**. Otherwise, is it English, including informal spellings (u, dat, lol, wtf)? → **EN**. *EN vs SG is one judgement about use in this message, not a checklist order.*
4. Chinese characters or pinyin? → **ZH**
5. Recognisable Malay → **MS**; recognisable Tamil → **TA**; recognisable Hokkien → **HK**
6. Can’t tell which non-English language → **UNK**

4.3 Rules where mistakes happen

- Informal English is still EN. “u” = you, “dat” = that, “alr” = already, “j” = just, “bef” = before. Internet slang (LOL, LMAO, IDK, WTF) is also **EN**. Don’t tag these **UNK** just because they’re abbreviated.
- “ya” / “yeah” / “ye” default to EN, not **SG** as they are informal English agreement. Only tag them **SG** if they are clearly working as a singlish particle in the message. If it just means “yeah/yep”, tag **EN**.
- **SG** covers particles and Singlish borrowings, by use. Particles (lah, leh, sia) and borrowed words (paiseh = embarrassed, kiasu = afraid to lose, sian = bored) all get **SG**. You’re not expected to know their origin or grammar so if it’s a colloquial Singlish item just label it **SG**.
- Romanised Chinese (pinyin) is ZH. “ni hao”, “xie xie”, “wo” → **ZH**. (But a romanised item that has become a conventional Singlish borrowing, e.g. an encouragement marker like *jiayou*, is **SG**. Judge by

use.)

- Laughter is OTH, even though it's letters. "HAHAHA", "hahaha", "HAHAHHAHA" → OTH.
- "ah" is usually SG when attached to a sentence ("wait ah", "okay ah"). Only a standalone English exclamation ("ah, I see") is EN. When in doubt → SG.
- The defaults for "ya" and "ah" point in opposite directions because in this data "ah" is nearly always a particle while "ya" is nearly always plain agreement. By default it should follow the more common use. Either way, clear context overrides it.
- Can't tell which Asian language? Best guess from context only use UNK if truly stuck. UNK is preferred here.

5 Harm labels (per message)

Answer the two questions independently.

Field	Labels
Q1 Harm type (the words)	<code>hate_speech · offensive · neither</code>
Q2 Aggression (the tone)	<code>genuine_aggression · mock_aggression · none</code>
Target	<code>self · addressee · third_party · group · none</code>
Confidence	<code>certain · uncertain</code>

When context may be used:

Task	Context allowed?
Language labels	No. tag each token by what it is
Q1 (harm type)	Minimal. only to interpret the current message's own content (e.g. is this word quoted? denied?)
Q2 (aggression) & target	Yes. relationship and preceding turns are how you tell mock from actual aggression

5.1 Q1: is there harmful or offensive content?

Judge only the linguistic content of the current message.

- `hate_speech`: attacks, demeans, excludes, or expresses hostility toward an individual or group because of a protected characteristic (race, religion, gender, sexuality, nationality, disability). E.g. "people like [group] shouldn't be here", or a demeaning generalisation about a protected group. In the Singaporean context, hostility framed around nationality, migration status, or ethnicity is a likely form to appear so watch for it specifically.
- `offensive`: explicit abusive or demeaning language: profanity, vulgarity, derogatory descriptions, or direct personal insults ("you useless", "you're pathetic", "wtf is this garbage"), but not targeting a protected group.
- `neither`: no swearing, no insult. (Most messages.)

Swear-word test

A swear word alone is not automatically offensive. If it's just venting the speaker's own feeling

(“ugh, damn”), that’s **neither**. (A message like “so sian” is **neither** too, but for a simpler reason: it contains no profanity or insult at all.) Take out the swear word: if the message is just “I’m annoyed/excited”, it’s **neither**. Mark **offensive** only when the profanity or insult has real bite: directed at a person, or used as an intentionally derogatory evaluation of something (“this movie is fucking garbage”).

Reported or quoted insults like “he called me stupid”, “she said wtf to me” describe harm but they don’t perform it → **neither**, unless the current speaker endorses or amplifies it (“he called you stupid and he’s right” → **offensive**).

5.2 Q1a: harmful content in non-English terms

Offensive and hateful language in this data will not always be in English, and if you only check for English profanity you will miss it. Romanised Hokkien and Malay vulgarities, in particular, are common in Singlish chat and would not be flagged by an English profanity list. Two rules:

- The same form/function rule applies in every language. A romanised Hokkien vulgarity used as pure venting (“knn the bus late again”) is the same case as English “ugh damn” → likely **neither** / **none**. The same vulgarity aimed at a person (“knn you”) has bite → **offensive**. Apply the swear-word test regardless of language.

Fluency fallback (same logic as UNK)

If you meet a non-English token you cannot confirm as profane or derogatory, don’t pass the message as clean, and don’t assume it’s harmful either. Label the message on the evidence you can understand: if the rest of the message and context are clean, that is **neither**. Set **confidence: uncertain**, and note the specific token for a fluent check at adjudication. The uncertain flag and the note are what pass the case on, don’t guess a content label, and don’t pick a direction you can’t justify. An unfamiliar token means you lack information, which is different from whether the message itself can be read two ways.

The one exception: if the surrounding message is itself clearly hostile (its structure, target, and cold context all point to an attack, and only the specific term is unconfirmable), label what the message is clearly doing (**offensive**, or **hate_speech** if it targets a protected characteristic), still mark **uncertain**, and note the token.

Mild illustrative examples (constructed): “you bodoh” (Malay *bodoh* = stupid), directed at the addressee → **offensive** · target **addressee**. A romanised Hokkien vulgarity in self-venting about the weather → **neither** · **none**.

5.3 Q2: is it aggressive, and is it real or playful?

Look at how the message is used.

- **genuine_aggression**: a real attack, threat, or anger, meant. E.g. “don’t ever message me again, I’m done with you.”
- **mock_aggression**: hostile-sounding but actually friendly: teasing or banter between people clearly close. Cues: laughing emoji, close-friend relationship, reciprocal teasing. E.g. “you idiot” with a laughing emoji.
- **none**: not aggressive at all.

To tell mock from genuine, use context and the relationship. Close friends joking, laughing emoji, mutual teasing → mock. Cold tone, real grievance, no joking → genuine. Can't tell even with context → best guess plus **confidence**: `uncertain` (Section 8.1).

5.4 Target

Target is who the harm or aggression is aimed at. Three rules:

- If offensive language is aimed at an object, product, situation, or event rather than a person, target is `none`.
- If Q1 = `neither` and Q2 = `none`, target is always `none`.
- If Q2 is `genuine_aggression` or `mock_aggression`, the target cannot be `none` and aggression is aimed at someone.

Self-directed messages: light self-deprecation venting frustration (“im so stupid lol”, “im damn useless sia” with laughter) works like venting → `neither` · `none` · target `none`. A harsh, serious self-attack with no playful cues → `offensive` · `none` · target `self`, and consider **confidence**: `uncertain`.

5.5 Harm decision tree

Q1 — harm type

1. Attacks a protected group (race / religion / gender / nationality / ...)? → `hate_speech`
2. Swearing, vulgarity, or a direct personal insult with bite (in *any* language)? → `offensive`
3. Otherwise → `neither`

Q2 — aggression (judged separately)

1. A real attack, threat, or anger, meant? → `genuine_aggression`
2. Hostile-sounding but friendly banter between close people? → `mock_aggression`
3. Otherwise → `none`
4. Can't tell mock vs genuine even with context → record your best guess and set **confidence**: `uncertain` (add a one-line note)

6 Worked examples (the hard cases)

Read all of these before you start labelling.

Ex 1 — Friendly teasing that sounds like a threat

Message	“U btr b”
Context	close friends joking about a race
Tokens	U·EN btr·EN b·EN (informal English, “you better be”)
Q1 type	<code>neither</code> : no swearing or insult
Q2 tone	<code>mock_aggression</code> : sounds threatening, but banter between close friends
Target	<code>addressee</code>
Confidence	<code>certain</code>

Why hostile-sounding words can be friendly; answer the two questions separately, so no harmful content but a playful-aggressive tone

Ex 2 — Swearing not aimed at a person

Message “but i heard they taste like shit”

Context friends discussing energy gels

Q1 type *offensive*: not because a swear word appears, but because it’s a derogatory evaluation of the gels (“taste like shit” has bite; contrast Ex 3)

Q2 tone *none*: nobody is attacked

Target *none*: the gels are a thing, not a person

Why the two questions are independent: offensive words, but no aggression toward anyone

Ex 3 — A swear word that’s just venting

Message “WAH DAMN SIAN”

Context complaining it started raining

Tokens WAH·SG DAMN·EN SIAN·SG (“sian” = fed-up)

Q1 type *neither*

Q2 tone *none*

Target *none*

Why take out “damn” and you get “wah, so bored”: pure feeling, no target. A swear word alone doesn’t make a message offensive (contrast Ex 2)

Ex 4 — Singlish particles and borrowings

Message “Aiya I paiseh sia”

Tokens Aiya·SG I·EN paiseh·SG sia·SG (“paiseh” = embarrassed)

Q1 type *neither*

Q2 tone *none*

Target *none*

Why an English “I” surrounded by Singlish items; tag each token on its own

Ex 5 — A swear-ish word that’s about a feeling

Message “i tried it bef but i j ended up being pissed while running”

Tokens all EN (“bef” = before, “j” = just)

Q1 type *neither*

Q2 tone *none*

Target *none*

Why “pissed” means the speaker *felt* annoyed and they describe their own emotion, not an attack. An angry-sounding word isn’t an aggressive act

Ex 6 — Mixed languages, one per token

Message	“睡觉time alr”
Tokens	睡觉·ZH time·EN alr·EN (“睡觉” = sleep, “alr” = already)
Q1 type	neither
Q2 tone	none
Why	a message can switch language word by word; tag each token in its own language

Ex 7 — Loudness is not aggression

Message	“NOT THIS AGAIN HAHAHHAHAHA”
Tokens	NOT·EN THIS·EN AGAIN·EN HAHAHHAHAHA·OTH
Q1 type	neither
Q2 tone	none
Target	none
Why	ALL CAPS and laughter read intense, but there’s no aggressive act here

Ex 8 — Negation (don’t flag what’s being denied)

Message	“he’s not stupid actually”
Tokens	all EN
Q1 type	neither
Q2 tone	none
Why	the word “stupid” appears, but it’s being denied so the message says the opposite of an insult. Read the whole message, not isolated words

Ex 9 — The uncertain case

Message	“you useless”
Context	close friends after one missed an easy goal
Tokens	you·EN useless·EN
Q1 type	offensive: a direct personal insult; no profanity needed
Q2 tone	mock_aggression: best guess: close friends, light context
Target	addressee
Confidence	uncertain: note: “could be genuine; context doesn’t settle it”
Why	Q1 is settled (“useless” is a real insult). What’s unclear is the <i>function</i> : when even context can’t tell mock from genuine, record your best guess and let the uncertain flag carry the ambiguity forward (Section 8.1)

Ex 10 — Genuine aggression (*constructed*)

Message	“don’t ever message me again. i’m done with you.”
Context	after a serious falling-out, cold tone
Tokens	all EN
Q1 type	neither: no profanity, slur, or explicit insult; genuine aggression needs none

Q2 tone	<code>genuine_aggression</code>
Target	<code>addressee</code>
Why	real aggression is cold, serious, and meant, with no swearing. Compare Ex 1: friendly despite threatening wording. Context, not vocabulary, tells them apart

Ex 11 — Hate speech (*constructed, rare here but central to the task*)

Message	“ugh [nationality] workers everywhere now, they shouldn’t even be allowed here”
Context	unprompted complaint, no joking
Q1 type	<code>hate_speech</code> : demeans and seeks to exclude a group because of nationality, a protected characteristic
Q2 tone	<code>genuine_aggression</code>
Target	<code>group</code>
Confidence	<code>certain</code>
Why	this is the line between <code>offensive</code> and <code>hate_speech</code> : attacking people for what they are is hate speech, even with no profanity. For an unconfirmable non-English slur, use the fluency fallback (Section 5.2)

Ex 12 — Self-directed talk (*constructed*)

Message	“im so stupid lol i forgot my keys again”
Context	light chat
Tokens	all <code>EN</code> , including lol (internet slang → <code>EN</code> , not <code>OTH</code>)
Q1 type	<code>neither</code>
Q2 tone	<code>none</code>
Target	<code>none</code>
Why	light self-deprecation with a playful cue (“lol”) works like venting. Compare a harsh, serious self-attack (“i’m worthless. i ruin everything.”) → <code>offensive · none · target self, uncertain</code> if the register is unclear

7 The order to annotate

Do every message in this fixed order so that it keeps everyone consistent:

1. Read the context (preceding turns plus relationship).
2. Language: tag each token.
3. Q1, harm type (`hate_speech` / `offensive` / `neither`).
4. Q2, aggression (`genuine` / `mock` / `none`).
5. Target (or `none`).
6. Confidence (`certain` / `uncertain`, add a one-line note if uncertain).

8 Workflow, disagreement & quality

8.1 Confidence: what “uncertain” means

Mark **uncertain** in either of two situations, and they are not the same:

- Insufficient evidence: the text or context doesn’t give you enough to decide (e.g. you don’t know who the speaker is, or a non-English token you can’t confirm). This kind often gets fixed by a clearer rule or more context.
- The other case is a message that really can be read two ways (Ex 9). No rule settles it, and the split is information.

Both get **confidence: uncertain**; the distinction matters at adjudication (below), so add a note saying which it is when it’s not obvious.

8.2 Pilot / calibration round

Before full independent annotation, run a **small calibration batch** (e.g. 15–25 turns) that everyone labels, then compare and talk through the splits. This is where unclear rules show up and get fixed before full annotation starts. Type-1 fixes (below) start here. Only after calibration do annotators move to fully independent labelling. No agreement scores from the calibration round are treated as final: it exists to refine the guidelines.

8.3 Independent annotation & adjudication

Three annotators per batch: one primary, two secondary.

1. Everyone labels independently. No discussion first: it skews the agreement numbers.
2. Labels are compared message by message: all agree → settled; two agree → majority recorded but the disagreement is **kept**; all differ → **flagged for adjudication**.
3. The primary leads a review of disagreements. The review looks for the cause of each disagreement.
4. Adjudication outcomes feed back into these guidelines. When a disagreement was caused by an unclear rule, update the guidelines (add a rule or an example) **before the next batch**, so the same confusion doesn’t recur.

Two kinds of disagreement could appear:

- Type 1, fixable (a rule was unclear). E.g. people split on “ya” because the rule was ambiguous. → Clarify the rule or add an example, so it won’t recur.
- Type 2, real (the message is genuinely ambiguous). E.g. “you useless”: readable as playful or cutting, and no rule can settle it. → The disagreement stays as it’s real information.

In review, ask: would a clearer rule have prevented this? Yes → Type 1, fix it. No → Type 2, keep it.

Genuinely ambiguous messages are not forced to a majority. We keep the split (e.g. 2 saw playful, 1 saw an edge) as the label. Forcing one confident label onto an ambiguous message manufactures certainty that isn’t there, and systems trained on those labels over-flag harmless chat.

8.4 Inter-annotator agreement

Report agreement with **Krippendorff's** α , per task (language, harm type, aggression). It works with three annotators and missing labels, and it doesn't require collapsing the kept splits into one gold label. Track α across batches, a calibration round and guideline updates should reduce the fixable disagreement while the irreducible share stays flat.

8.5 Quality expectations

- Use **uncertain** freely; it's the correct answer when a message can't be called confidently.
- Read the context every time before harm labels.
- Answer the two questions separately, every message.
- When torn between two labels, record your best guess and mark **uncertain**, with a one-line note. Don't jump to the harsher label on a hunch. Don't pick the milder one "to be safe" either: softening hides the disagreement, and the uncertain flag is what carries it forward.
- Don't infer demographics. Don't assume anyone's race, gender, religion, etc. unless explicitly stated.
- Add a one-line note on hard calls.